

Machine Learning Engineer Nanodegree

Capstone Project: Facial Keypoints Detection

Daniel Vargas

December 15, 2019

I. Definition

Project Overview

Facial Keypoints (facial landmarks) detection is an important and challenging problem in the field of **computer vision**, which involves detecting facial keypoints like centers and corners of **eyes**, **nose**, and **mouth**, etc. The problem is to predict the **(x, y)** real-valued coordinates in the space of image pixels of the facial keypoints for a given face image.

Facial features vary greatly from one individual to another, and even for a single individual there is a large amount of variation due to **pose**, **size**, **position**, etc. The problem becomes even more challenging when the face images are taken under different **illumination conditions**, **viewing angles**, etc.

Solving this problem that can provide the building blocks for several applications, such as:

- tracking faces in images and video
- analysing facial expressions
- detecting dysmorphic facial signs for medical diagnosis
- biometrics / face recognition

In the past few years, advancements in facial keypoints detection have been made by implementing ****Deep Convolutional Neural Networks (DCNN)****.

Relevant academic research on this domain can be found in

- [Facial Keypoints Detection](#)
- [Facial Key Points Detection using Deep Convolutional Neural Network - NaimishNet](#).

I chose this specific challenge because I currently work in the medical diagnosis field. I expect this project to help me understand facial keypoints recognition in a deeper way.

Datasets and Inputs

The data was acquired from the [Facial Keypoints Detection Kaggle](#) competition.

Data files

- **training.csv**: list of **7049 training images**. Each row contains the (x, y) coordinates for 15 keypoints, and image data as row-ordered list of pixels.
- **test.csv**: list of **1783 test images**. Each row contains ImageId and image data as row-ordered list of pixels

Problem Statement

The objective of this project is to accurately predict the facial keypoints (facial landmarks) of a face image. My hypothesis is, that this prediction can be performed based on a training set containing accurate facial keypoints, through a regression approach.

A **Convolutional Neural Network (CNN)** will be applied to predict the facial keypoints. A **CNN** was chosen for this problem because:

- This is a computer vision problem that requires capturing features for prediction
- CNNs are very useful in capturing features in images
- The expected responses (coordinates) make this a regression problem

A simple **Multilayer Perceptron (MLP)** will be used as a baseline model for comparison.

Metrics

The metric used to measure performance of the model is **Root Mean Squared Error (RMSE)**:

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - y_i)^2}$$

RMSE is very common and is a suitable general-purpose error metric in regression problems. Compared to the **Mean Absolute Error**, **RMSE** punishes large errors.

Network Strategy

- Data augmentation will be included if results are not satisfactory
- The network's architecture is as follows:
 - Input layer
 - Convolution layers
 - Max Pooling layers
 - Batch Normalization layers
 - Fully Connected layers
 - Dropout layers
 - Prediction layer

II. Analysis

Data Exploration

Each predicted keypoint is specified by an (x,y) real-valued pair in the space of pixel indices. There are 15 keypoints, which represent the following elements of the face:

left_eye_center	right_eye_center	right_eye_center
left_eye_inner_corner	left_eye_outer_corner	right_eye_inner_corner

left_eyebrow_inner_end	left_eyebrow_outer_end	right_eyebrow_inner_end
right_eyebrow_outer_end	nose_tip	mouth_left_corner
mouth_right_corner	mouth_center_top_lip	mouth_center_bottom_lip

- *Left and right here refers to the point of view of the subject*
- *In some examples, some of the target keypoint positions are missing (encoded as missing entries in the csv, i.e., with nothing between two commas)*
- *The input image is given in the last field of the data files, and consists of a list of pixels (ordered by row), as integers in (0,255). The images are 96x96 pixels*

Exploratory Visualization

The data is summarized as follows.

Data columns (total 31 columns):


```

left_eye_center_x          7039 non-null float64
left_eye_center_y          7039 non-null float64
right_eye_center_x         7036 non-null float64
right_eye_center_y         7036 non-null float64
left_eye_inner_corner_x    2271 non-null float64
left_eye_inner_corner_y    2271 non-null float64
left_eye_outer_corner_x    2267 non-null float64
left_eye_outer_corner_y    2267 non-null float64
right_eye_inner_corner_x   2268 non-null float64
right_eye_inner_corner_y   2268 non-null float64
right_eye_outer_corner_x   2268 non-null float64
right_eye_outer_corner_y   2268 non-null float64
left_eyebrow_inner_end_x   2270 non-null float64
left_eyebrow_inner_end_y   2270 non-null float64
left_eyebrow_outer_end_x   2225 non-null float64
left_eyebrow_outer_end_y   2225 non-null float64
right_eyebrow_inner_end_x  2270 non-null float64
right_eyebrow_inner_end_y  2270 non-null float64
right_eyebrow_outer_end_x  2236 non-null float64
right_eyebrow_outer_end_y  2236 non-null float64
nose_tip_x                 7049 non-null float64
nose_tip_y                 7049 non-null float64
mouth_left_corner_x        2269 non-null float64
mouth_left_corner_y        2269 non-null float64
mouth_right_corner_x       2270 non-null float64
mouth_right_corner_y       2270 non-null float64
mouth_center_top_lip_x     2275 non-null float64
mouth_center_top_lip_y     2275 non-null float64
mouth_center_bottom_lip_x  7016 non-null float64
mouth_center_bottom_lip_y  7016 non-null float64
Image                     7049 non-null object

```

```
dtypes: float64(30), object(1)
```

Also, below is an example group of images.



Algorithms and Techniques

For this problem, I will use a **Convolutional Neural Network (CNN)**: In this, will use a Sequential (there's only a single input) model With 3 different **Conv2D** layers each having a **max pooling** layer having a **pool size** and **stride** of (2, 2). Each layer I have also added **batch normalization** and **dropouts** to mitigate overfitting. At the end, I have also added 3 **fully connected** layers With **dropouts**. In this, have used **adam** optimizer having **epochs** set to 100 and a **batch size** of 128.

This is the model's summary:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 96, 96, 16)	80
dropout_2 (Dropout)	(None, 96, 96, 16)	0

max_pooling2d_1 (MaxPooling2)	(None, 48, 48, 16)	0
batch_normalization_1 (Batch Normalization)	(None, 48, 48, 16)	64
conv2d_2 (Conv2D)	(None, 44, 44, 32)	12832
max_pooling2d_2 (MaxPooling2)	(None, 22, 22, 32)	0
dropout_3 (Dropout)	(None, 22, 22, 32)	0
batch_normalization_2 (Batch Normalization)	(None, 22, 22, 32)	128
conv2d_3 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_3 (MaxPooling2)	(None, 9, 9, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 9, 9, 64)	256
conv2d_4 (Conv2D)	(None, 7, 7, 128)	73856
max_pooling2d_4 (MaxPooling2)	(None, 3, 3, 128)	0
dropout_4 (Dropout)	(None, 3, 3, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 3, 3, 128)	512
flatten_1 (Flatten)	(None, 1152)	0
dense_4 (Dense)	(None, 500)	576500
dropout_5 (Dropout)	(None, 500)	0
dense_5 (Dense)	(None, 128)	64128
dropout_6 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 30)	3870

=====

Total params: 783,490
 Trainable params: 783,010
 Non-trainable params: 480

Benchmark

The benchmark Multilayer Perceptron (MLP): In this, will use a `sequential` model With 3 different layers followed by an activation function `Relu`, and will also add a `dropout` after the first layer. I have used `SGD` optimizer for this using 50 `epochs` and a `batch size` of 128.

This is the model's summary:

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	2359552
activation_1 (Activation)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
activation_2 (Activation)	(None, 128)	0
dense_3 (Dense)	(None, 30)	3870
Total params: 2,396,318		
Trainable params: 2,396,318		
Non-trainable params: 0		

III. Methodology

Data Preprocessing

Various operations were performed on the data for training.

- Convert the image values to numpy arrays (The **Image** column has pixel values separated by spaces)
- Drop all rows that have missing values in them
- Scale all pixel values to (0, 1) (normalize)
- Scale target coordinates to (-1, 1)
- Shuffle train data to mitigate overfitting

Implementation

The implemented models are baseline model (**MLP.h5**) and the final model (**CNN.h5**).


IV. Results

Model Evaluation and Validation

The final model was evaluated with the dataset provided in **Kaggle** and tested by submitting the results to the competition.

Model	Private Score	Public Score
Baseline model (MLP.h5)	4.24427	4.28892
Final model (CNN.h5)	3.97870	3.99583

Below, there's a screenshot of the submissions.



Detect the location of keypoints on face images

175 teams · 3 years ago

OverviewDataNotebooksDiscussionLeaderboardRulesTeamMy SubmissionsLate Submission

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
udacity-MLP.csv	just now	3 seconds	0 seconds	4.28892

Complete

[Jump to your position on the leaderboard](#)


You may select up to 5 submissions to be used to count towards your final leaderboard score. If 5 submissions are not selected, they will be automatically chosen based on your best submission scores on the public leaderboard. In the event that automatic selection is not suitable, manual selection instructions will be provided in the competition rules or by official forum announcement.

Your final score may not be based on the same exact subset of data as the public leaderboard, but rather a different private data subset of your full submission — your public score is only a rough indication of what your final score is.

You should thus choose submissions that will most likely be best overall, and not necessarily on the public subset.

>_

kaggle competitions submit -c facial-keypoints-detection -f submission.csv -m "Message"

 ?

7 submissions for DanielVargas

Sort byMost recent

AllSuccessfulSelected

Submission and Description	Private Score	Public Score	Use for Final Score
udacity-MLP.csv just now by DanielVargas udacity benchmark model MLP	4.24427	4.28892	<input type="checkbox"/>
udacity-final-CNN.csv 28 minutes ago by DanielVargas udacity final model CNN	3.97870	3.99583	<input type="checkbox"/>

Justification

Based on the score results from **Kaggle**, the final model (**CNN.h5**) performed better than the benchmark model (**MLP.h5**) on the test set.

V. Conclusion

Reflection

The following steps were taken to complete this process:

- 1. Downloaded the dataset from kaggle
- 2. Performed data preprocessing
- 3. Trained a baseline model for comparison (**MLP.h5**)
- 4. Trained a final model for implementation (**CNN.h5**)

5. Converted the results to `.csv` and submitted them to score them on **Kaggle**
6. Chose the best model based on the models' individual scores

I learned a lot from this project:

- The importance on investing the time to prepare the data the right way to have a smooth training
- The importance of having in mind that the model might be implemented, taking into account not only metrics, but also prediction performance

Improvement

There are many ways in which the final model can be improved. The trade-offs of these improvements would depend on the final purpose of the model.

- Perform hyperparameters optimization (e.g. random search, bayes optimization)
- Perform random image augmentation on the training set (e.g. rotations, translations, zoom-in, zoom-out, blur, etc.)
- Quantize the final model before conversion to reduce size and prediction speed